

```
In [1]: import os
import numpy as np
import pandas as pd
from sklearn import model_selection
from sklearn.metrics import roc_curve
from sklearn.metrics import auc, confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.ensemble import AdaBoostClassifier
from sklearn.grid_search import GridSearchCV
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\cross_validation.py:41: DeprecationWarning: This module was deprecated in version 0.18 in favor of the model_selection module into which all the refactored classes and functions are moved. Also note that the interface of the new CV iterators are different from that of this module. This module will be removed in 0.20.

"This module will be removed in 0.20.", DeprecationWarning)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\grid_search.py:42: DeprecationWarning: This module was deprecated in version 0.18 in favor of the model_selection module into which all the refactored classes and functions are moved. This module will be removed in 0.20.

DeprecationWarning)

```
In [2]: os.chdir("D:/Great Lakes PGPDSE/Great Lakes/13 Ensemble Techniques/Mini Project")
```

Reading the dataset

```
In [4]: hr = pd.read_csv("HR_Employee_Attrition_Dat.csv")
```

Converting Yes / No values in Attrition column to 1 / 0

```
In [5]: cleanup_nums = {"Attrition": {"Yes": 1, "No": 0}}
```


```
In [6]: hr.replace(cleanup_nums, inplace=True)
```

In [25]: `hr.head()`

Out[25]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Educational
0	41	1	Travel_Rarely	1102	Sales	1	2
1	49	0	Travel_Frequently	279	Research & Development	8	1
2	37	1	Travel_Rarely	1373	Research & Development	2	2
3	33	0	Travel_Frequently	1392	Research & Development	3	4
4	27	0	Travel_Rarely	591	Research & Development	2	1

5 rows × 35 columns



Splitting to training and testing data

In [26]: `from sklearn.cross_validation import train_test_split
import random`

In [27]: `np.random.seed(40)
train,test = train_test_split(hr, test_size = 0.3)`

Splitting target variable and independent variables

```
In [28]: X1 = train[['Age', 'BusinessTravel', 'DailyRate', 'Department',
    'DistanceFromHome', 'Education', 'EducationField',
    'EnvironmentSatisfaction', 'Gender', 'HourlyRate', 'JobInvolvement',
    'JobLevel', 'JobRole', 'JobSatisfaction', 'MaritalStatus',
    'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked', 'OverTime',
    'PercentSalaryHike', 'PerformanceRating', 'RelationshipSatisfaction',
    'StockOptionLevel', 'TotalWorkingYears', 'TrainingTimesLastYear',
    'WorkLifeBalance', 'YearsAtCompany', 'YearsInCurrentRole',
    'YearsSinceLastPromotion', 'YearsWithCurrManager']]
X2 = test[['Age', 'BusinessTravel', 'DailyRate', 'Department',
    'DistanceFromHome', 'Education', 'EducationField',
    'EnvironmentSatisfaction', 'Gender', 'HourlyRate', 'JobInvolvement',
    'JobLevel', 'JobRole', 'JobSatisfaction', 'MaritalStatus',
    'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked', 'OverTime',
    'PercentSalaryHike', 'PerformanceRating', 'RelationshipSatisfaction',
    'StockOptionLevel', 'TotalWorkingYears', 'TrainingTimesLastYear',
    'WorkLifeBalance', 'YearsAtCompany', 'YearsInCurrentRole',
    'YearsSinceLastPromotion', 'YearsWithCurrManager']]
```

```
In [29]: y_train = train["Attrition"]
y_test = test["Attrition"]
```

Categorical Variable to Numerical Variables

```
In [30]: X_train = pd.get_dummies(X1)
X_test = pd.get_dummies(X2)
X_train.columns
```

```
Out[30]: Index(['Age', 'DailyRate', 'DistanceFromHome', 'Education',
'EnvironmentSatisfaction', 'HourlyRate', 'JobInvolvement', 'JobLevel',
'JobSatisfaction', 'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorke
d',
'PercentSalaryHike', 'PerformanceRating', 'RelationshipSatisfaction',
'StockOptionLevel', 'TotalWorkingYears', 'TrainingTimesLastYear',
'WorkLifeBalance', 'YearsAtCompany', 'YearsInCurrentRole',
'YearsSinceLastPromotion', 'YearsWithCurrManager',
'BusinessTravel_Non-Travel', 'BusinessTravel_Travel_Frequently',
'BusinessTravel_Travel_Rarely', 'Department_Human Resources',
'Department_Research & Development', 'Department_Sales',
'EducationField_Human Resources', 'EducationField_Life Sciences',
'EducationField_Marketing', 'EducationField_Medical',
'EducationField_Other', 'EducationField_Technical Degree',
'Gender_Female', 'Gender_Male', 'JobRole_Healthcare Representative',
'JobRole_Human Resources', 'JobRole_Laboratory Technician',
'JobRole_Manager', 'JobRole_Manufacturing Director',
'JobRole_Research Director', 'JobRole_Research Scientist',
'JobRole_Sales Executive', 'JobRole_Sales Representative',
'MaritalStatus_Divorced', 'MaritalStatus_Married',
'MaritalStatus_Single', 'OverTime_No', 'OverTime_Yes'],
dtype='object')
```

Fitting AdaBoost model

```
In [31]: model = AdaBoostClassifier(n_estimators=30)
model.fit(X_train, y_train)
```

```
Out[31]: AdaBoostClassifier(algorithm='SAMME.R', base_estimator=None,
learning_rate=1.0, n_estimators=30, random_state=None)
```

```
In [32]: pred_y_train = model.predict(X_train)
pred_y_train
```

```
Out[32]: array([0, 0, 0, ..., 0, 0, 1], dtype=int64)
```

Classification accuracy of the model in train data

```
In [33]: score = accuracy_score(y_train, pred_y_train)
         score
```

Out[33]: 0.8974732750242954

AUC of the model in train data

```
In [34]: y_train_prob = model.predict_proba(X_train)
         fpr, tpr, thresholds = roc_curve(y_train, y_train_prob[:,1])
         auc(fpr, tpr)
```

Out[34]: 0.9006297150478626

Classification accuracy of the model in test data

```
In [35]: pred_y_test = model.predict(X_test)
pred_y_test
```

[illegible]

```
In [36]: score_test = accuracy_score(y_test, pred_y_test)
          score_test
```

Out[36]: 0.8968253968253969

AUC of the model in train data

```
In [37]: y_test_prob = model.predict_proba(X_test)
         fpr, tpr, thresholds = roc_curve(y_test, y_test_prob[:,1])
         auc(fpr, tpr)
```

```
Out[37]: 0.8435021178282633
```

Here the AUC for test data differs alot compared to that of train data. That is the model is overfitting. So we need to do some parameter tuning

Cross validatin

```
In [38]: scores = model_selection.cross_val_score(model, X_train, y_train, cv = 10, scoring='roc_auc')
         scores.mean()
```

```
Out[38]: 0.8504723144558511
```

```
In [21]: scores.std()
```

```
Out[21]: 0.04768426938729329
```

So by cross validation we get the correct AUC of the model, that is 85.04% is the correct AUC with standard deviation of 0.047. For test data it is almost same, but auc for training data is high. so the model is overfitted. Need to tune the model.

Checking important variables

```
In [39]: import pandas as pd
feature_imp = pd.Series(model.feature_importances_, index=X_train.columns).sort_
_values(ascending=False)
feature_imp
```

```

Out[39]: MonthlyIncome      0.100000
        DailyRate           0.066667
        JobInvolvement       0.066667
        JobSatisfaction       0.066667
        NumCompaniesWorked    0.066667
        YearsWithCurrManager   0.066667
        OverTime_Yes          0.033333
        WorkLifeBalance       0.033333
        BusinessTravel_Travel_Frequently 0.033333
        BusinessTravel_Non-Travel 0.033333
        YearsSinceLastPromotion 0.033333
        YearsAtCompany         0.033333
        StockOptionLevel       0.033333
        TrainingTimesLastYear  0.033333
        TotalWorkingYears      0.033333
        Department_Research & Development 0.033333
        RelationshipSatisfaction 0.033333
        EnvironmentSatisfaction 0.033333
        DistanceFromHome       0.033333
        OverTime_No           0.033333
        Age                   0.033333
        EducationField_Technical Degree 0.033333
        JobRole_Laboratory Technician 0.033333
        EducationField_Human Resources 0.000000
        PerformanceRating      0.000000
        MaritalStatus_Single   0.000000
        MaritalStatus_Married  0.000000
        Education              0.000000
        MaritalStatus_Divorced 0.000000
        HourlyRate             0.000000
        JobRole_Sales Representative 0.000000
        JobLevel               0.000000
        JobRole_Sales Executive 0.000000
        JobRole_Research Scientist 0.000000
        MonthlyRate            0.000000
        JobRole_Research Director 0.000000
        PercentSalaryHike      0.000000
        JobRole_Manufacturing Director 0.000000
        Department_Sales       0.000000
        JobRole_Manager        0.000000
        JobRole_Human Resources 0.000000
        JobRole_Healthcare Representative 0.000000
        Gender_Male            0.000000
        YearsInCurrentRole     0.000000
        Gender_Female          0.000000
        EducationField_Other   0.000000
        EducationField_Medical 0.000000
        EducationField_Marketing 0.000000
        Department_Human Resources 0.000000
        EducationField_Life Sciences 0.000000
        BusinessTravel_Travel_Rarely 0.000000
        dtype: float64

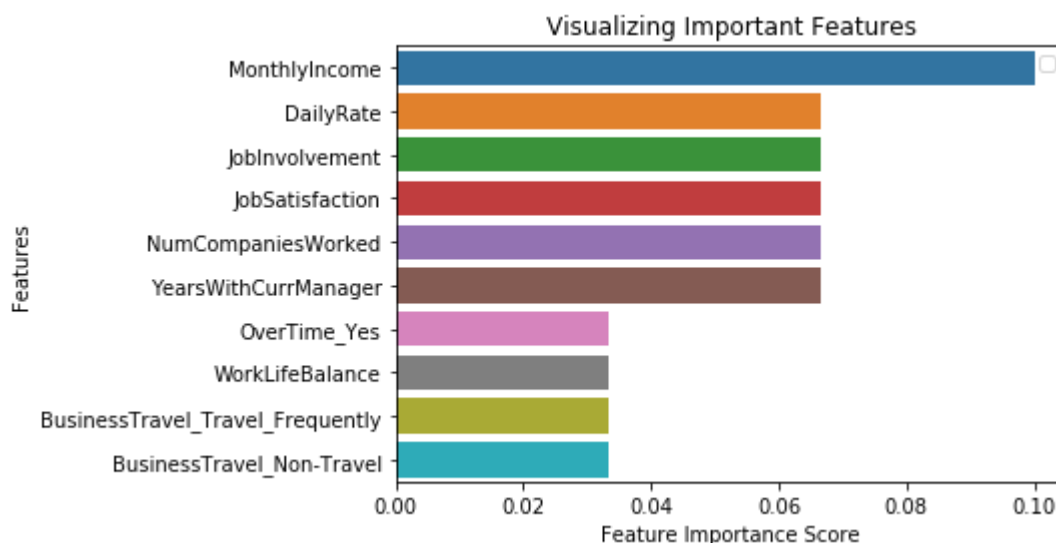
```



```
In [40]: import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
# Creating a bar plot
feature_imp=feature_imp[0:10,]

sns.barplot(x=feature_imp, y=feature_imp.index)
# Add labels to your graph
plt.xlabel('Feature Importance Score')
plt.ylabel('Features')
plt.title("Visualizing Important Features")
plt.legend()
plt.show()
```

No handles with labels found to put in legend.



Here most important variable is monthly income followed by daily rate, jobInvolvement etc

Parameter Tuning

```
In [43]: param_dist = {"n_estimators": np.arange(10, 20),
                        "learning_rate": [0.1, 0.2, 0.3, 0.5, 0.6, 0.7, 0.8, 0.9, 1],
                        }
```

```
In [44]: tree = AdaBoostClassifier(random_state=None)
tree_cv = GridSearchCV(tree, cv = 10,
                        param_grid=param_dist,
                        n_jobs = 3)
tree_cv.fit(X_train, y_train)

Out[44]: GridSearchCV(cv=10, error_score='raise',
                      estimator=AdaBoostClassifier(algorithm='SAMME.R', base_estimator=None,
                                                    learning_rate=1.0, n_estimators=50, random_state=None),
                      fit_params={}, iid=True, n_jobs=3,
                      param_grid={'n_estimators': array([10, 11, 12, 13, 14, 15, 16, 17, 18,
19]), 'learning_rate': [0.1, 0.2, 0.3, 0.5, 0.6, 0.7, 0.8, 0.9, 1]},
                      pre_dispatch='2*n_jobs', refit=True, scoring=None, verbose=0)
```

Building the model using best combination of parameters

```
In [45]: print("Tuned Decision Tree parameter : {}".format(tree_cv.best_params_))
classifier = tree_cv.best_estimator_
classifier.fit(X_train,y_train)

Tuned Decision Tree parameter : {'learning_rate': 0.9, 'n_estimators': 15}

Out[45]: AdaBoostClassifier(algorithm='SAMME.R', base_estimator=None,
                             learning_rate=0.9, n_estimators=15, random_state=None)
```

Checking AUC of the tuned model

For train data

```
In [46]: y_train_prob = classifier.predict_proba(X_train)
fpr, tpr, thresholds = roc_curve(y_train, y_train_prob[:,1])
auc_d = auc(fpr, tpr)
auc_d

Out[46]: 0.8702762688080915
```

For test data

```
In [47]: y_test_prob = classifier.predict_proba(X_test)
fpr, tpr, thresholds = roc_curve(y_test, y_test_prob[:,1])
auc_h = auc(fpr, tpr)
auc_h

Out[47]: 0.8084809395456295
```

Scoring and Rank ordering

```
In [48]: Prediction = classifier.predict_proba(X_train)
train["prob_score"] = Prediction[:,1]
```

C:\ProgramData\Anaconda3\lib\site-packages\ipykernel_launcher.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

decile code

```
In [49]: def deciles(x):
decile = pd.Series(index=[0,1,2,3,4,5,6,7,8,9])
for i in np.arange(0.1,1.1,0.1):
    decile[int(i*10)]=x.quantile(i)
def z(x):
    if x<decile[1]: return(1)
    elif x<decile[2]: return(2)
    elif x<decile[3]: return(3)
    elif x<decile[4]: return(4)
    elif x<decile[5]: return(5)
    elif x<decile[6]: return(6)
    elif x<decile[7]: return(7)
    elif x<decile[8]: return(8)
    elif x<decile[9]: return(9)
    elif x<=decile[10]: return(10)
    else: return(np.NaN)
s=x.map(z)
return(s)
```

```
In [51]: def Rank_Ordering(X,y,Target):
X['decile']=deciles(X[y])
Rank=X.groupby('decile').apply(lambda x: pd.Series([
    np.min(x[y]),
    np.max(x[y]),
    np.mean(x[y]),
    np.size(x[y]),
    np.sum(x[Target]),
    np.size(x[Target][x[Target]==0]),
]),
    index=["min_resp", "max_resp", "avg_resp",
           "cnt", "cnt_resp", "cnt_non_resp"])
    ).reset_index()
Rank = Rank.sort_values(by='decile',ascending=False)
Rank["rrate"] = Rank["cnt_resp"]*100/Rank["cnt"]
Rank["cum_resp"] = np.cumsum(Rank["cnt_resp"])
Rank["cum_non_resp"] = np.cumsum(Rank["cnt_non_resp"])
Rank["cum_resp_pct"] = Rank["cum_resp"]/np.sum(Rank["cnt_resp"])
Rank["cum_non_resp_pct"] = Rank["cum_non_resp"]/np.sum(Rank["cnt_non_resp"])
Rank["KS"] = Rank["cum_resp_pct"] - Rank["cum_non_resp_pct"]
Rank
return(Rank)
```

Rank ordering of train data

```
In [52]: Rank = Rank_Ordering(train,"prob_score","Attrition")
Rank
```

C:\ProgramData\Anaconda3\lib\site-packages\ipykernel_launcher.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

Out[52]:

	decile	min_resp	max_resp	avg_resp	cnt	cnt_resp	cnt_non_resp	rrate	curr
9	10	0.492985	0.552292	0.508124	207.0	149.0	58.0	71.980676	149.
8	9	0.478632	0.492753	0.485165	205.0	73.0	132.0	35.609756	222.
7	8	0.470529	0.478534	0.474212	206.0	50.0	156.0	24.271845	272.
6	7	0.462778	0.470368	0.466580	209.0	16.0	193.0	7.655502	288.
5	6	0.456808	0.462664	0.460072	203.0	12.0	191.0	5.911330	300.
4	5	0.450552	0.456662	0.453513	213.0	11.0	202.0	5.164319	311.
3	4	0.443132	0.450531	0.447345	200.0	11.0	189.0	5.500000	322.
2	3	0.435527	0.443102	0.440030	200.0	5.0	195.0	2.500000	327.
1	2	0.424101	0.435035	0.430296	212.0	5.0	207.0	2.358491	332.
0	1	0.389532	0.423488	0.413146	203.0	2.0	201.0	0.985222	334.

Here Best KS=0.6136 and response rate in 10th decile = 72.68%

Rank ordering for test data

```
In [53]: Prediction_h = classifier.predict_proba(X_test)
test["prob_score"] = Prediction_h[:,1]

Rank_h = Rank_Ordering(test,"prob_score","Attrition")
Rank_h
```

C:\ProgramData\Anaconda3\lib\site-packages\ipykernel_launcher.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

Out[53]:

	decile	min_resp	max_resp	avg_resp	cnt	cnt_resp	cnt_non_resp	rrate	cum_
9	10	0.491897	0.552292	0.508638	89.0	57.0	32.0	64.044944	57.0
8	9	0.479137	0.491757	0.484948	88.0	24.0	64.0	27.272727	81.0
7	8	0.471644	0.479030	0.474969	88.0	21.0	67.0	23.863636	102.0
6	7	0.463812	0.471616	0.467919	90.0	8.0	82.0	8.888889	110.0
5	6	0.458810	0.463751	0.461099	86.0	6.0	80.0	6.976744	116.0
4	5	0.451048	0.458749	0.454577	88.0	5.0	83.0	5.681818	121.0
3	4	0.443638	0.451027	0.447738	88.0	7.0	81.0	7.954545	128.0
2	3	0.437537	0.443423	0.440701	90.0	6.0	84.0	6.666667	134.0
1	2	0.428106	0.437471	0.431890	87.0	4.0	83.0	4.597701	138.0
0	1	0.389532	0.427822	0.416982	88.0	2.0	86.0	2.272727	140.0

Here on testing the model on test data we are getting KS value of 0.5088 and response rate of 64.04% in the top decile. This looks very good but compared to the training data this is less. The model looks bit over fitted.

Here the employees who all comes in the 10th decile are most likely to leave the company. Let's assume that training of new employee costs 1000 dollar and since we know which employee is likely to leave next month, and propose him/her a bonus program worth 500 to keep him for next 6 months, we are 500 dollar to keep him for next 6 months, we are 500 dollar on plus and keep experienced, well-trained employee under the hood, with higher morale. So better to take care of these employees and company should be prepared to find substitution for those.