

```
In [2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import os
%matplotlib inline
```

Setting the working Directory

```
In [21]: os.chdir("D:/Great Lakes PGPDSE/Great Lakes/13 Ensemble Techniques/Mini Project")
```

Reading the data set

```
In [22]: hr = pd.read_csv("HR_Employee_Attrition_Dat.csv")
hr.head()
```

Out[22]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Educational
0	41	Yes	Travel_Rarely	1102	Sales	1	2
1	49	No	Travel_Frequently	279	Research & Development	8	1
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2
3	33	No	Travel_Frequently	1392	Research & Development	3	4
4	27	No	Travel_Rarely	591	Research & Development	2	1

5 rows × 35 columns



```
In [23]: hr.shape
```

Out[23]: (2940, 35)

There are 35 columns and 2940 data entries in the file

```
In [24]: hr.dtypes
```

```
Out[24]: Age                int64
Attrition                  object
BusinessTravel             object
DailyRate                 int64
Department                 object
DistanceFromHome          int64
Education                 int64
EducationField             object
EmployeeCount             int64
EmployeeNumber            int64
EnvironmentSatisfaction   int64
Gender                    object
HourlyRate                int64
JobInvolvement            int64
JobLevel                  int64
JobRole                   object
JobSatisfaction           int64
MaritalStatus             object
MonthlyIncome             int64
MonthlyRate               int64
NumCompaniesWorked        int64
Over18                    object
OverTime                  object
PercentSalaryHike         int64
PerformanceRating         int64
RelationshipSatisfaction  int64
StandardHours             int64
StockOptionLevel          int64
TotalWorkingYears         int64
TrainingTimesLastYear     int64
WorkLifeBalance           int64
YearsAtCompany            int64
YearsInCurrentRole        int64
YearsSinceLastPromotion   int64
YearsWithCurrManager      int64
dtype: object
```


Checking summary statistics

In [25]: `hr.describe()`

Out[25]:

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	Em
count	2940.000000	2940.000000	2940.000000	2940.000000	2940.0	294
mean	36.923810	802.485714	9.192517	2.912925	1.0	147
std	9.133819	403.440447	8.105485	1.023991	0.0	848
min	18.000000	102.000000	1.000000	1.000000	1.0	1.0
25%	30.000000	465.000000	2.000000	2.000000	1.0	735
50%	36.000000	802.000000	7.000000	3.000000	1.0	147
75%	43.000000	1157.000000	14.000000	4.000000	1.0	220
max	60.000000	1499.000000	29.000000	5.000000	1.0	294

8 rows × 26 columns




Checking for missing values

In [26]: `hr[hr.isnull().any(axis=1)]`

Out[26]:

Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education
-----	-----------	----------------	-----------	------------	------------------	-----------

0 rows × 35 columns



There are no missing values

```
In [27]: hr.nunique()
```

```
Out[27]: Age                43
Attrition                2
BusinessTravel           3
DailyRate              886
Department              3
DistanceFromHome        29
Education               5
EducationField           6
EmployeeCount            1
EmployeeNumber         2940
EnvironmentSatisfaction  4
Gender                  2
HourlyRate              71
JobInvolvement           4
JobLevel                5
JobRole                  9
JobSatisfaction          4
MaritalStatus            3
MonthlyIncome          1349
MonthlyRate             1427
NumCompaniesWorked       10
Over18                   1
OverTime                 2
PercentSalaryHike        15
PerformanceRating        2
RelationshipSatisfaction  4
StandardHours            1
StockOptionLevel         4
TotalWorkingYears        40
TrainingTimesLastYear     7
WorkLifeBalance          4
YearsAtCompany           37
YearsInCurrentRole        19
YearsSinceLastPromotion   16
YearsWithCurrManager      18
dtype: int64
```

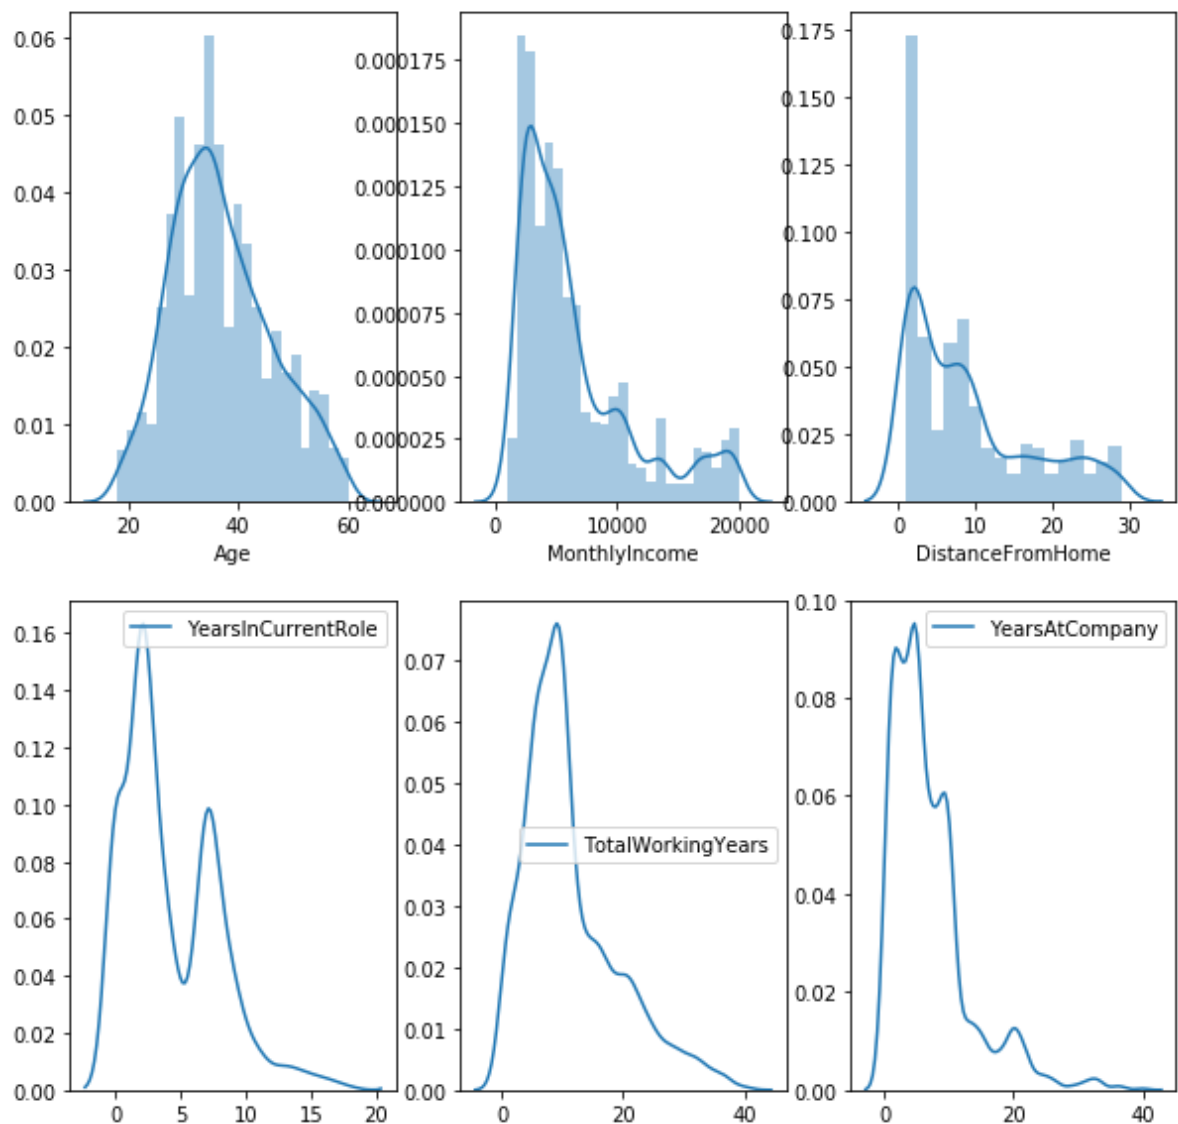
Here the value for columns, Over18, StandardHours and EmployeeCount same for all rows, we can eliminate these columns.

```
In [28]: del hr["Over18"]
del hr["EmployeeCount"]
del hr["StandardHours"]
```

Plotting Histogram for various factors in the same plot space.

```
In [29]: fig,ax = plt.subplots(2,3, figsize=(10,10))
plt.suptitle("Distribution of various factors", fontsize=20)
sns.distplot(hr['Age'], ax = ax[0,0])
sns.distplot(hr['MonthlyIncome'], ax = ax[0,1])
sns.distplot(hr['DistanceFromHome'], ax = ax[0,2])
sns.kdeplot(hr['YearsInCurrentRole'], ax = ax[1,0])
sns.kdeplot(hr['TotalWorkingYears'], ax = ax[1,1])
sns.kdeplot(hr['YearsAtCompany'], ax = ax[1,2])
plt.show()
```

Distribution of various factors

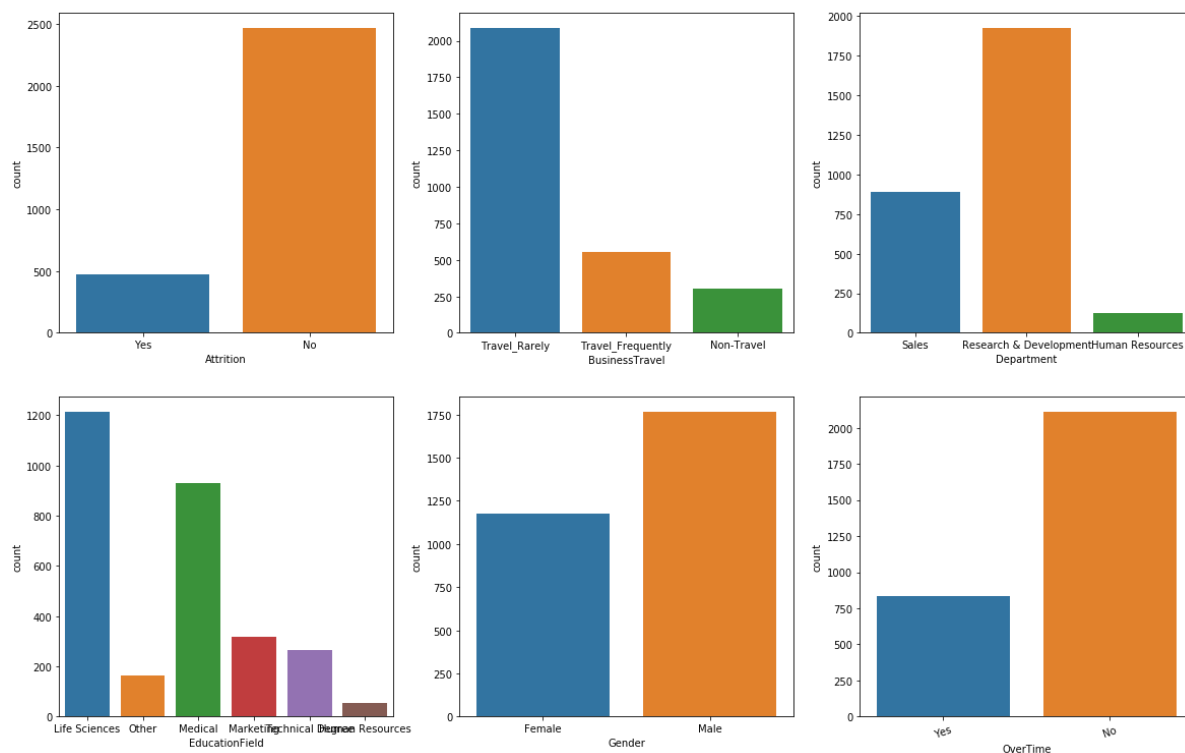


Only Age is showing the normal distribution in the data , all other variables not showing the normal distribution

Checking the distribution of various factors

```
In [30]: fig,ax = plt.subplots(2,3, figsize=(20,20))
plt.suptitle("Distribution of various factors", fontsize=20)
sns.countplot(hr['Attrition'], ax = ax[0,0])
sns.countplot(hr['BusinessTravel'], ax = ax[0,1])
sns.countplot(hr['Department'], ax = ax[0,2])
sns.countplot(hr['EducationField'], ax = ax[1,0])
sns.countplot(hr['Gender'], ax = ax[1,1])
sns.countplot(hr['OverTime'], ax = ax[1,2])
plt.xticks(rotation=20)
plt.subplots_adjust(bottom=0.4)
plt.show()
```

Distribution of various factors




In the data we have, Attrition data is not distributed equally in this dataset. So this dataset is unbalanced data set.

Checking gender wise distribution


```
In [31]: fig,ax = plt.subplots(2,3, figsize=(20,20))
plt.suptitle("Distribution of various factors", fontsize=20)
sns.barplot(hr['Gender'],hr['DistanceFromHome'],hue = hr['Attrition'], ax = ax
[0,0]);
sns.barplot(hr['Gender'],hr['YearsAtCompany'],hue = hr['Attrition'], ax = ax[0
,1]);
sns.barplot(hr['Gender'],hr['TotalWorkingYears'],hue = hr['Attrition'], ax = a
x[0,2]);
sns.barplot(hr['Gender'],hr['YearsInCurrentRole'],hue = hr['Attrition'], ax =
ax[1,0]);
sns.barplot(hr['Gender'],hr['YearsSinceLastPromotion'],hue = hr['Attrition'],
ax = ax[1,1]);
sns.barplot(hr['Gender'],hr['NumCompaniesWorked'],hue = hr['Attrition'], ax =
ax[1,2]);
plt.show()
```

Distribution of various factors



Plotting a correlation map for all numeric variables

```
In [32]: f,ax = plt.subplots(figsize=(18, 18))
sns.heatmap(hr.corr(), annot=True, linewidths=.5, fmt= '.1f',ax=ax)
plt.show()
```



Converting Yes / No values in Attrition column to 1 / 0

```
In [34]: cleanup_nums = {"Attrition": {"Yes": 1, "No": 0}}
```



```
In [35]: hr.replace(cleanup_nums, inplace=True)
hr.head()
```

Out[35]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Educational
0	41	1	Travel_Rarely	1102	Sales	1	2
1	49	0	Travel_Frequently	279	Research & Development	8	1
2	37	1	Travel_Rarely	1373	Research & Development	2	2
3	33	0	Travel_Frequently	1392	Research & Development	3	4
4	27	0	Travel_Rarely	591	Research & Development	2	1

5 rows × 32 columns



Splitting to training and testing data

```
In [36]: from sklearn.cross_validation import train_test_split
import random
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\cross_validation.py:41: DeprecationWarning: This module was deprecated in version 0.18 in favor of the model_selection module into which all the refactored classes and functions are moved. Also note that the interface of the new CV iterators are different from that of this module. This module will be removed in 0.20.
 "This module will be removed in 0.20.", DeprecationWarning)

```
In [38]: np.random.seed(1010)
train,test = train_test_split( hr, test_size = 0.3)
```

Checking distribution of Attribution in each group

```
In [39]: hr.Attrition.value_counts()
```

```
Out[39]: 0    2466
         1     474
         Name: Attrition, dtype: int64
```

```
In [14]: 474/(2466+474)
```

```
Out[14]: 0.16122448979591836
```

```
In [40]: test.Attrition.value_counts()
```

```
Out[40]: 0    727
         1    155
         Name: Attrition, dtype: int64
```

```
In [43]: 155/(155+727)
```

```
Out[43]: 0.17573696145124718
```

```
In [41]: train.Attrition.value_counts()
```

```
Out[41]: 0    1739
         1     319
         Name: Attrition, dtype: int64
```

```
In [44]: 319/(319+1739)
```

```
Out[44]: 0.15500485908649175
```

Split of attribution is approsxame for both train and test data

Splitting target variable y_train and y_test and independent variables X1 and X2 variables

```
In [45]: X1 = train[['Age', 'BusinessTravel', 'DailyRate', 'Department',
                    'DistanceFromHome', 'Education', 'EducationField',
                    'EnvironmentSatisfaction', 'Gender', 'HourlyRate', 'JobInvolvement',
                    'JobLevel', 'JobRole', 'JobSatisfaction', 'MaritalStatus',
                    'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked', 'OverTime',
                    'PercentSalaryHike', 'PerformanceRating', 'RelationshipSatisfaction',
                    'StockOptionLevel', 'TotalWorkingYears', 'TrainingTimesLastYear',
                    'WorkLifeBalance', 'YearsAtCompany', 'YearsInCurrentRole',
                    'YearsSinceLastPromotion', 'YearsWithCurrManager']]
        X2 = test[['Age', 'BusinessTravel', 'DailyRate', 'Department',
                   'DistanceFromHome', 'Education', 'EducationField',
                   'EnvironmentSatisfaction', 'Gender', 'HourlyRate', 'JobInvolvement',
                   'JobLevel', 'JobRole', 'JobSatisfaction', 'MaritalStatus',
                   'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked', 'OverTime',
                   'PercentSalaryHike', 'PerformanceRating', 'RelationshipSatisfaction',
                   'StockOptionLevel', 'TotalWorkingYears', 'TrainingTimesLastYear',
                   'WorkLifeBalance', 'YearsAtCompany', 'YearsInCurrentRole',
                   'YearsSinceLastPromotion', 'YearsWithCurrManager']]
```

```
In [46]: y_train = train["Attrition"]
         y_test = test["Attrition"]
```

Categorical Variable to Numerical Variables

```
In [47]: X_train = pd.get_dummies(X1)
X_test = pd.get_dummies(X2)
X_train.columns
```

```
Out[47]: Index(['Age', 'DailyRate', 'DistanceFromHome', 'Education',
'EnvironmentSatisfaction', 'HourlyRate', 'JobInvolvement', 'JobLevel',
'JobSatisfaction', 'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked',
'PercentSalaryHike', 'PerformanceRating', 'RelationshipSatisfaction',
'StockOptionLevel', 'TotalWorkingYears', 'TrainingTimesLastYear',
'WorkLifeBalance', 'YearsAtCompany', 'YearsInCurrentRole',
'YearsSinceLastPromotion', 'YearsWithCurrManager',
'BusinessTravel_Non-Travel', 'BusinessTravel_Travel_Frequently',
'BusinessTravel_Travel_Rarely', 'Department_Human Resources',
'Department_Research & Development', 'Department_Sales',
'EducationField_Human Resources', 'EducationField_Life Sciences',
'EducationField_Marketing', 'EducationField_Medical',
'EducationField_Other', 'EducationField_Technical Degree',
'Gender_Female', 'Gender_Male', 'JobRole_Healthcare Representative',
'JobRole_Human Resources', 'JobRole_Laboratory Technician',
'JobRole_Manager', 'JobRole_Manufacturing Director',
'JobRole_Research Director', 'JobRole_Research Scientist',
'JobRole_Sales Executive', 'JobRole_Sales Representative',
'MaritalStatus_Divorced', 'MaritalStatus_Married',
'MaritalStatus_Single', 'OverTime_No', 'OverTime_Yes'],
dtype='object')
```

Applying Random Forest for the above data

Scaling the data

```
In [48]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

```
In [49]: from sklearn.ensemble import RandomForestClassifier
```

```
In [50]: model = RandomForestClassifier()
```

```
In [51]: model.fit(X_train,y_train)
```

```
Out[51]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
max_depth=None, max_features='auto', max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=1,
oob_score=False, random_state=None, verbose=0,
warm_start=False)
```

```
In [52]: pred_y_train = model.predict(X_train)
pred_y_train
```

```
Out[52]: array([0, 0, 0, ..., 0, 0, 0], dtype=int64)
```

Let us see the classification accuracy of our model

```
In [54]: from sklearn.metrics import accuracy_score
score = accuracy_score(y_train, pred_y_train)
score
```

```
Out[54]: 0.9941690962099126
```

AUC

```
In [55]: from sklearn.metrics import roc_curve
from sklearn.metrics import auc, confusion_matrix

y_train_prob = model.predict_proba(X_train)
fpr, tpr, thresholds = roc_curve(y_train, y_train_prob[:,1])
auc(fpr, tpr)
```

```
Out[55]: 0.9999296969216265
```

Checking for Test data

```
In [56]: pred_y_test = model.predict(X_test)
pred_y_test

## Let us see the classification accuracy of our model
score_test = accuracy_score(y_test, pred_y_test)
score_test
```

```
Out[56]: 0.9319727891156463
```

```
In [57]: y_test_prob = model.predict_proba(X_test)
fpr, tpr, thresholds = roc_curve(y_test, y_test_prob[:,1])
auc(fpr, tpr)
```

```
Out[57]: 0.9500377157563118
```

```
In [58]: from sklearn import model_selection
scores = model_selection.cross_val_score(model, X_train, y_train, cv = 10, scoring='roc_auc')
scores.mean()
```

```
Out[58]: 0.9422174385315734
```

```
In [59]: scores.std()
```

```
Out[59]: 0.020981436334551497
```

So by cross validation we get the correct AUC of the model, that is 94.26% is the correct AUC with standard deviation of 0.020. For test data and model selection through cross validation is almost same, but auc for training data is high. so the model is overfitted.

Checking important variables and arranging in the descending order

```

In [61]: import pandas as pd
feature_imp = pd.Series(model.feature_importances_, index= ['Age', 'DailyRate',
'DistanceFromHome', 'Education',
'EnvironmentSatisfaction', 'HourlyRate', 'JobInvolvement', 'JobLevel',
'JobSatisfaction', 'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked'
,
'PercentSalaryHike', 'PerformanceRating', 'RelationshipSatisfaction',
'StockOptionLevel', 'TotalWorkingYears', 'TrainingTimesLastYear',
'WorkLifeBalance', 'YearsAtCompany', 'YearsInCurrentRole',
'YearsSinceLastPromotion', 'YearsWithCurrManager',
'BusinessTravel_Non-Travel', 'BusinessTravel_Travel_Frequently',
'BusinessTravel_Travel_Rarely', 'Department_Human Resources',
'Department_Research & Development', 'Department_Sales',
'EducationField_Human Resources', 'EducationField_Life Sciences',
'EducationField_Marketing', 'EducationField_Medical',
'EducationField_Other', 'EducationField_Technical Degree',
'Gender_Female', 'Gender_Male', 'JobRole_Healthcare Representative',
'JobRole_Human Resources', 'JobRole_Laboratory Technician',
'JobRole_Manager', 'JobRole_Manufacturing Director',
'JobRole_Research Director', 'JobRole_Research Scientist',
'JobRole_Sales Executive', 'JobRole_Sales Representative',
'MaritalStatus_Divorced', 'MaritalStatus_Married',
'MaritalStatus_Single', 'OverTime_No', 'OverTime_Yes'] ).sort_values(as
cending=False)
feature_imp

```

```


Out[61]: MonthlyIncome      0.082483
         Age                 0.064031
         DailyRate          0.052367
         DistanceFromHome   0.048545
         HourlyRate         0.046515
         TotalWorkingYears   0.041863
         NumCompaniesWorked  0.041041
         YearsAtCompany      0.038896
         OverTime_Yes       0.035306
         MonthlyRate        0.034621
         EnvironmentSatisfaction 0.032266
         YearsSinceLastPromotion 0.031774
         YearsWithCurrManager 0.030587
         WorkLifeBalance     0.030251
         StockOptionLevel    0.029649
         PercentSalaryHike   0.029339
         YearsInCurrentRole  0.028878
         TrainingTimesLastYear 0.024228
         JobSatisfaction     0.024214
         RelationshipSatisfaction 0.022147
         JobInvolvement      0.021971
         MaritalStatus_Single 0.017762
         Education           0.016408
         JobLevel            0.015559
         BusinessTravel_Travel_Frequently 0.014292
         OverTime_No         0.013723
         EducationField_Medical 0.010010
         JobRole_Sales Representative 0.009031
         Gender_Male         0.008289
         JobRole_Laboratory Technician 0.008053
         EducationField_Marketing 0.007692
         BusinessTravel_Travel_Rarely 0.007366
         Department_Sales    0.006842
         Gender_Female       0.006837
         Department_Research & Development 0.006449
         EducationField_Life Sciences 0.006368
         EducationField_Technical Degree 0.005867
         MaritalStatus_Married 0.005426
         JobRole_Research Scientist 0.005175
         BusinessTravel_Non-Travel 0.005132
         EducationField_Other 0.004881
         MaritalStatus_Divorced 0.004880
         JobRole_Sales Executive 0.004716
         PerformanceRating   0.004259
         JobRole_Manufacturing Director 0.002619
         Department_Human Resources 0.002533
         EducationField_Human Resources 0.002338
         JobRole_Healthcare Representative 0.002277
         JobRole_Research Director 0.002218
         JobRole_Human Resources 0.001650
         JobRole_Manager     0.000375
         dtype: float64

```

```
In [62]: import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
# Creating a bar plot
feature_imp=feature_imp[0:10,]

sns.barplot(x=feature_imp, y=feature_imp.index)
# Add labels to your graph
plt.xlabel('Feature Importance Score of given variable')
plt.ylabel('Features')
plt.title("Visualizing Important Features")
plt.legend()
plt.show()
```

No handles with labels found to put in legend.



Here most important variable is MonthlyIncome followed by Age and Daily rate

Parameter Tuning for the Random forest model

```
In [63]: import time
from sklearn.grid_search import GridSearchCV
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\grid_search.py:42: DeprecationWarning: This module was deprecated in version 0.18 in favor of the model_selection module into which all the refactored classes and functions are moved. This module will be removed in 0.20.
DeprecationWarning)


```

In [64]: np.random.seed(42)
         start = time.time()

         param_dist = {

             'n_estimators': np.arange(10, 20),
             'max_depth': [2, 3, 4],
             'bootstrap': [True, False],
             'max_features': ['auto', 'sqrt', 'log2', None],
             'criterion': ['gini', 'entropy']}

         cv_rf = GridSearchCV(model, cv = 10,
                               param_grid=param_dist,
                               n_jobs = 3)

         cv_rf.fit(X_train, y_train)
         print('Best Parameters using grid search: \n',
               cv_rf.best_params_)
         end = time.time()
         print('Time taken in grid search: {0: .2f}'.format(end - start))

```

```

Best Parameters using grid search:
{'bootstrap': True, 'criterion': 'gini', 'max_depth': 4, 'max_features': None,
'n_estimators': 19}
Time taken in grid search: 236.21

```

```

In [65]: classifier = cv_rf.best_estimator_
         classifier.fit(X_train, y_train)

```

```

Out[65]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                                max_depth=4, max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=19, n_jobs=1,
                                oob_score=False, random_state=None, verbose=0,
                                warm_start=False)

```

AUC

```

In [66]: y_train_prob = classifier.predict_proba(X_train)
         fpr, tpr, thresholds = roc_curve(y_train, y_train_prob[:,1])
         auc_d = auc(fpr, tpr)
         auc_d

```

```

Out[66]: 0.8709199428201629

```

```

In [67]: y_test_prob = classifier.predict_proba(X_test)
         fpr, tpr, thresholds = roc_curve(y_test, y_test_prob[:,1])
         auc_h = auc(fpr, tpr)
         auc_h

```

```

Out[67]: 0.8300927363890491

```

Now the model seems more good. AUC for both data is almost same with 4% difference between train dataset and test dataset.

```
In [69]: Prediction = classifier.predict_proba(X_train)
train["prob_score"] = Prediction[:,1]
```

C:\ProgramData\Anaconda3\lib\site-packages\ipykernel_launcher.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

scoring and rank ordering

decile code

```
In [71]: def deciles(x):
decile = pd.Series(index=[0,1,2,3,4,5,6,7,8,9])
for i in np.arange(0.1,1.1,0.1):
    decile[int(i*10)]=x.quantile(i)
def z(x):
    if x<decile[1]: return(1)
    elif x<decile[2]: return(2)
    elif x<decile[3]: return(3)
    elif x<decile[4]: return(4)
    elif x<decile[5]: return(5)
    elif x<decile[6]: return(6)
    elif x<decile[7]: return(7)
    elif x<decile[8]: return(8)
    elif x<decile[9]: return(9)
    elif x<=decile[10]: return(10)
    else: return(np.NaN)
s=x.map(z)
return(s)
```

```
In [72]: def Rank_Ordering(X,y,Target):
X['decile']=deciles(X[y])
Rank=X.groupby('decile').apply(lambda x: pd.Series([
    np.min(x[y]),
    np.max(x[y]),
    np.mean(x[y]),
    np.size(x[y]),
    np.sum(x[Target]),
    np.size(x[Target][x[Target]==0]),
]),
    index=["min_resp", "max_resp", "avg_resp",
           "cnt", "cnt_resp", "cnt_non_resp"])
    ).reset_index()
Rank = Rank.sort_values(by='decile',ascending=False)
Rank["rrate"] = Rank["cnt_resp"]*100/Rank["cnt"]
Rank["cum_resp"] = np.cumsum(Rank["cnt_resp"])
Rank["cum_non_resp"] = np.cumsum(Rank["cnt_non_resp"])
Rank["cum_resp_pct"] = Rank["cum_resp"]/np.sum(Rank["cnt_resp"])
Rank["cum_non_resp_pct"]=Rank["cum_non_resp"]/np.sum(Rank["cnt_non_resp"])
Rank["KS"] = Rank["cum_resp_pct"] - Rank["cum_non_resp_pct"]
Rank
return(Rank)
```

Rank ordering for train data

```
In [74]: Rank = Rank_Ordering(train,"prob_score","Attrition")
Rank
```

C:\ProgramData\Anaconda3\lib\site-packages\ipykernel_launcher.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

Out[74]:

	decile	min_resp	max_resp	avg_resp	cnt	cnt_resp	cnt_non_resp	rrate	curr
9	10	0.297405	0.950011	0.522539	206.0	161.0	45.0	78.155340	161.
8	9	0.191513	0.295821	0.232902	207.0	63.0	144.0	30.434783	224.
7	8	0.145209	0.190643	0.166641	205.0	32.0	173.0	15.609756	256.
6	7	0.118834	0.145195	0.132132	205.0	21.0	184.0	10.243902	277.
5	6	0.101833	0.118789	0.109218	206.0	8.0	198.0	3.883495	285.
4	5	0.092359	0.101824	0.096665	206.0	8.0	198.0	3.883495	293.
3	4	0.078050	0.092211	0.087245	206.0	8.0	198.0	3.883495	301.
2	3	0.059630	0.078050	0.067872	211.0	6.0	205.0	2.843602	307.
1	2	0.054601	0.059538	0.056443	203.0	5.0	198.0	2.463054	312.
0	1	0.047246	0.054419	0.050603	203.0	7.0	196.0	3.448276	319.

Response rate in top deciles is 80.25% and Highest ks value 59.94%. This looks bit overfitting.

Rank ordering for test data

```
In [75]: Prediction_h = classifier.predict_proba(X_test)
test["prob_score"] = Prediction_h[:,1]

Rank_h = Rank_Ordering(test,"prob_score","Attrition")
Rank_h
```

C:\ProgramData\Anaconda3\lib\site-packages\ipykernel_launcher.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

Out[75]:

	decile	min_resp	max_resp	avg_resp	cnt	cnt_resp	cnt_non_resp	rrate	cum_
9	10	0.301090	0.936781	0.507944	89.0	64.0	25.0	71.910112	64.0
8	9	0.184474	0.300052	0.224480	88.0	28.0	60.0	31.818182	92.0
7	8	0.148623	0.183026	0.165121	89.0	18.0	71.0	20.224719	110.0
6	7	0.123403	0.147540	0.135241	87.0	13.0	74.0	14.942529	123.0
5	6	0.103882	0.123064	0.111437	92.0	15.0	77.0	16.304348	138.0
4	5	0.095047	0.103834	0.099146	84.0	4.0	80.0	4.761905	142.0
3	4	0.085524	0.094750	0.090676	89.0	3.0	86.0	3.370787	145.0
2	3	0.061764	0.085384	0.070817	87.0	1.0	86.0	1.149425	146.0
1	2	0.054601	0.061516	0.057357	92.0	6.0	86.0	6.521739	152.0
0	1	0.049909	0.054419	0.050659	85.0	3.0	82.0	3.529412	155.0

Here on testing the model on test data we are getting KS value of 49.50% and response rate of 70.22% in the top decile. This looks very good but compared to the training data this is less. The model looks bit over fitted.

Here the employees who all comes in the 10th,9th and 8th decile are most likely to leave the company. Let's assume that training of new employee costs 1000 dollar and since we know which employee is likely to leave next month, and propose him/her a bonus program worth 500 to keep him for next 6 months, we are 500 dollar to keep him for next 6 months,we are 500 dollar on plus and keep experienced, well-trained employee under the hood, with higher morale. So better to take care of these employees and company should be prepared to find substitution for those.Also to find the better experienced resource company will get time to search better alternative resource.